

LLVM Pass and Code Instrumentation

Prof. Moonzoo Kim,
School of Computing, KAIST

Motivating Example

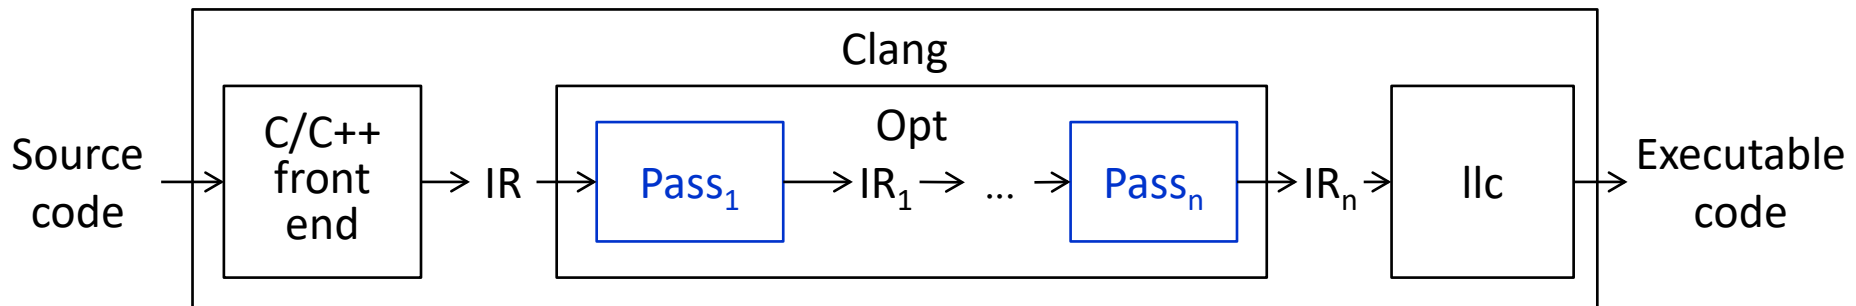
- (30 pts) Find bugs in the following program that has multiple bugs
 - Write down buggy lines, bugs, and explain the bugs as many as possible
 - Write down a code instrumentor using Clang which inserts `assert()` to report runtime failures due to the bugs you detected.
- For example, to report a div-by-zero crash, your code should insert `assert(z!=0)` immediately before `x=y/z`;

```
//example1.c
#include <malloc.h>
#include <stdio.h>
#include <string.h>
void f() {
    char* mem = NULL;
    int length;
    char buf[100];
    // file descriptor 0 is connected to keyboard
    read(0, &length, sizeof(int));
    int r=read(0, &buf,length>100 ?
        100:length);
    mem = malloc(r + 1);
    buf[r] = 0;
    strcpy(mem, buf);
    printf(mem);
    fflush(stdout);
}
```

Which tool do you prefer for the task? Clang? LLVM IR?

Pass in LLVM

- A Pass receives an LLVM IR and performs analyses and/or transformations.
 - Using `opt`, it is possible to run each Pass.
- A Pass can be executed in a middle of compiling process from source code to binary code.
 - The pipeline of Passes is arranged by Pass Manager



LLVM Pass Framework

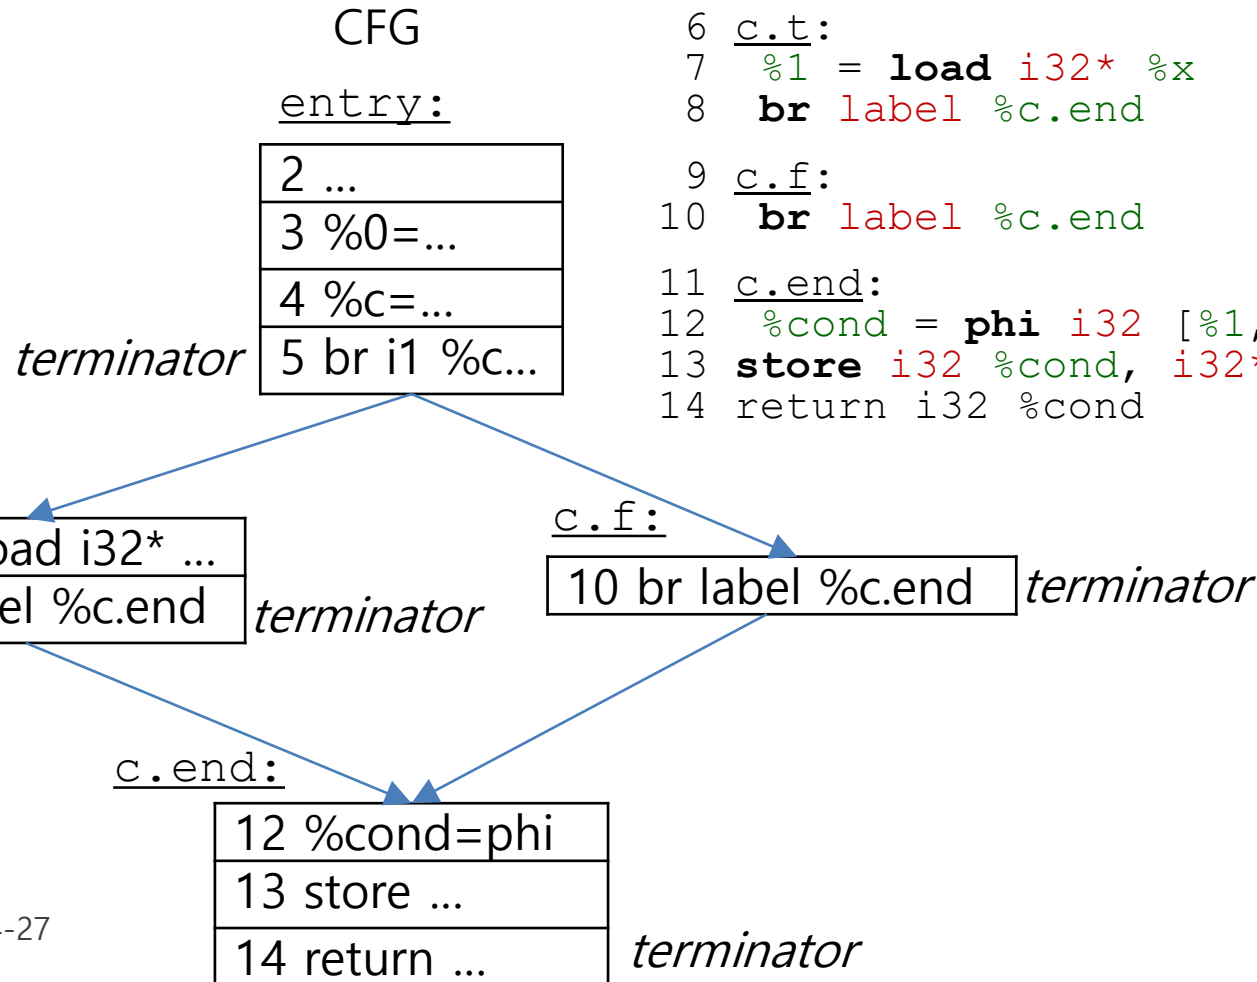
- The LLVM Pass Framework is the library to manipulate an AST of LLVM IR (<http://llvm.org/doxygen/index.html>)
- An LLVM Pass is an implementation of a subclass of the Pass class
 - Each Pass is defined as visitor on a certain type of LLVM AST nodes
 - There are six subclasses of Pass
 - ModulePass: visit each module (file)
 - CallGraphSCCPass: visit each set of functions with caller-call relations in a module (useful to draw a call graph)
 - FunctionPass: visit each function in a module
 - LoopPass: visit each set of basic blocks of a loop in each function
 - RegionPass: visit the basic blocks not in any loop in each function
 - BasicBlockPass: visit each basic block in each function

Control Flow Graph (CFG) at LLVM IR

```
int f() {
    int y;
    y = (x > 0) ? x : 0 ;
    return y;
}
```




```
1 entry:
2 ...
3 %0 = load i32* %x
4 %c = icmp sgt i32 %0 0
5 br i1 %c, label %c.t, %c.f
6 c.t:
7 %1 = load i32* %x
8 br label %c.end
9 c.f:
10 br label %c.end
11 c.end:
12 %cond = phi i32 [%1,%c.t],[0,%c.f]
13 store i32 %cond, i32* %y
14 return i32 %cond
```



Example Pass

- Let's create *IntWrite* that aim to monitor all history of 32-bit integer variable updates (definitions)
 - Implemented as a `FunctionPass`
 - Produces a text file where it records which variable is defined as which value at which code location.
- *IntWrite* instruments a target program to insert a *probe* before every integer writing operation, which extracts runtime information

```
01 int f(int x) {  
...  
10     y = x ;  
11     z = y + x ;
```



```
void _probe_(int l,char *fn,int v){ ...  
    fprintf(fp,"Store %d in %s @line %d\n",  
            v,fn,l) ;}  
  
...  
_probe_(10, "f", x) ;  
10 y = x ;  
_probe_(11, "f", y+x) ;  
11 z = y + x ;
```

Module Class

- A `Module` instance stores all information related to the LLVM IR created by a target program file (functions, global variables, etc.)
- APIs (public methods)
 - `getModuleIdentifier()`: return the name of the module
 - `getFunction(StringRef Name)`: return the `Function` instance whose identifier is `Name` in the module
 - `getOrInsertFunction(StringRef Name, Type *ReturnType, ...)`: add a new `Function` instance whose identifier is `Name` to the module
 - `getGlobalVariable(StringRef Name)`: return the `GlobalVariable` instance whose identifier is `Name` in the module

Type Class

- A `Type` instance is used for representing the data type of registers, variables, and function arguments.
- Static members
 - `Type::getVoidTy(...)` : void type
 - `Type::getInt8Ty(...)` : 8-bit unsigned integer (char) type
 - `Type::getInt32Ty(...)` : 32-bit unsigned integer type
 - `Type::getInt8PtrTy(...)` : 8-bit pointer type
 - `Type::getDoubleTy(...)` : 64-bit IEEE floating pointer type

FunctionPass Class (1/2)

- `FunctionPass::doInitialization(Module &)`
 - Executed once for a module (file) before any visitor method execution
 - Do necessary initializations, and modify the given `Module` instances (e.g., add a new function declaration)
- `FunctionPass::doFinalization(Module &)`
 - Executed once for a module (file) before after all visitor method executions
 - Export the information obtained from the analysis or the transformation, any wrap-up

Example

- IntWrite should insert a new function `_init_` at the beginning of the target program's main function
 - `_init_()` is to open an output file

```
01 virtual bool doInitialization(Module & M) {
02     if(M.getFunction(StartingRef("_init_")) != NULL) {
03         errs() << "_init_() already exists." ;
04         exit(1) ;
05     }                                     check if _init_() already exists

06     FunctionType *fty =
07         FunctionType::get(Type::getVoidTy(M.getContext()), false) ;
08     fp_init_ = M.getOrInsertFunction("_init_", fty) ;
09     ...                                 add a new declaration _init_()
10     return true ;
11 }
```

FunctionPass Class (2/2)

- `runOnFunction(Function &)`
 - Executed once for every function defined in the module
 - Read and modify the target function definition
- `Function Class`
 - `getFunctionType()`: returns the `FunctionType` instance that contains the information on the types of function arguments.
 - `getEntryBlock()`: returns the `BasicBlock` instance of the entry basic block.
 - `begin()`: the head of the `BasicBlock` iterator
 - `end()`: the end of the `BasicBlock` iterator

Example

```
01  virtual bool runOnFunction(Function &F) {
02      cout << "Analyzing " << F->getName() << "\n" ;
03      for (Function::iterator i = F.begin(); i != F.end(); i++){
04          runOnBasicBlock(*i) ;
05      }
06      return true; //You should return true if F was modified. False otherwise.
07  }
```

BasicBlock Class

- A `BasicBlock` instance contains a list of instructions
- APIs
 - `begin()`: return the iterator of the beginning of the basic block
 - `end()`: return the iterator of the end of the basic block
 - `getFirstInsertionPt()`: return the first iterator (i.e., the first instruction location) where a new instruction can be added safely (i.e., after phi instruction and debug intrinsic)
 - `getTerminator()`: return the terminator instruction
 - `splitBasicBlock(iterator I, ...)`: split the basic block into two at the instruction of `I` by inserting an unconditional jump

Instruction Class

- An `Instruction` instance contains the information of an LLVM IR instruction.
- Each type of instruction has a subclass of `Instruction` (e.g. `LoadInst`, `BranchInst`)
- APIs
 - `getOpcode()`: returns the opcode which indicates the instruction type
 - `getOperand(unsigned i)`: return the i-th operand
 - `getDebugLoc()`: obtain the debugging data that contains the information on the corresponding code location
 - `isTerminator()`, `isBinaryOp()`, `isCast()`, ...

Example

```
01 bool runOnBasicBlock(BasicBlock &B) {
02     for(BasicBlock::iterator i = B.begin(); i != B.end(); i++){
03         if(i->getOpcode() == Instruction::Store &&
04             i->getOperand(0)->getType() == Type::getInt32Ty(ctx)){
05             StoreInst * st = dyn_cast<StoreInst>(i);
06             int loc = st->getDebugLoc().getLine(); //code location
07             Value * var = st->getPointerOperand(); //variable
08             Value * val = st->getOperand(0); // value
09             /* insert a function call */
10         }
11     }
12     return true ;
13 }
```

How to Insert New Instructions

- `IRBuilder` class provides a uniform API for inserting instructions to a basic block.
 - `IRBuilder(Instruction *p)`: create an `IRBuilder` instance that can insert instructions **right before** `Instruction *p`
- APIs
 - `CreateAdd(Value *LHS, Value *RHS, ...)`: create an add instruction whose operands are `LHS` and `RHS` at the predefined location, and then returns the `Value` instance of the target operand
 - `CreateCall(Value *Callee, Value *Arg, ...)`: add a new call instruction to function `Callee` with the argument as `Arg`
 - `CreateSub()`, `CreateMul()`, `CreateAnd()`, ...

Value Class

- A `Value` is a super class of all entities in LLVM IR such as a constant, a register, a variable, and a function.
- The register defined by an `Instruction` is represented as a `Value` instance.
- APIs
 - `getType()`: returns the `Type` instance of a `Value` instance.
 - `getName()`: return the name from the source code.

Example

```
IRBuilder<> * IRB;

00 if(i->getOpcode() == Instruction::Store &&
01     i->getOperand(0)->getType() == Type::getInt32Ty(ctx) {

02     StoreInst * st = dyn_cast<StoreInst>(i);

03     int loc = st->getDebugLoc().getLine(); //code location
04     Value * val = st->getOperand(0); // target register
05
06     IRB->SetInsertPoint(&(*i));
07     Value * args[3] ;
08     args[0] = ConstantInt::get(intTy, loc, false) ;
09     args[1] = IRB->CreateGlobalStringPtr(funcname,"");
10     args[2] = val ;
11     IRB->CreateCall(p_probe, args, Twine(""));
    // p_probe should be created before by using
    // getOrInsertFunction() and target code should be compiled
    // with the function definition pointed by p_probe.
    // See IntWrite.cpp and IntWrite.c which contains the
    // definition of probe function

12 }
```



intwrite-pass.cpp

LLVM Fuction pass to monitor writing integer values at runtime
- Pass is compiled as a shared library and stored in the LLVM library directory



intwrite-rt.c

Runtime module that will be used by the pass
- Compiled as an object file and stored in the LLVM library directory

```
// intwrite-rt.c runtime module
```

```
#include <stdio.h>
#include <stdlib.h>
```

```
FILE * fp ;
```

```
extern void _final_() { fclose(fp) ;}
```

```
extern void _init_() { fp = fopen("log", "w") ; atexit(_final_) ;}
```

```
extern void _probe_(int line, char *func, int val) {
    if(line!=-1) fprintf(fp, "Store value %d in Function %s, line %d\n",val,func,line) ;
    else fprintf(fp, "Store value %d in unknown location\n", val) ;
}
```



intwrite-pass.cpp

LLVM Function pass to monitor writing integer values at runtime
- Pass is compiled as a shared library and stored in the LLVM library directory

```
// intwrite-pass.cpp LLVM pass
```

```
class IntWrite: public FunctionPass { ... }
```

```
/* The code in the remaining part is to register this Pass to  
 * LLVM Pass Manager such that LLVM/Clang can use it. */
```

```
char IntWrite::ID = 0;
```

```
static RegisterPass<IntWrite> X("IntWrite", "IntWrite Pass", false , false);
```

```
static void registerPass(const PassManagerBuilder &, legacy::PassManagerBase &PM) {  
    PM.add(new IntWrite());  
}
```

```
static RegisterStandardPasses  
    RegisterPassOpt(PassManagerBuilder::EP_ModuleOptimizerEarly, registerPass);
```

```
static RegisterStandardPasses  
    RegisterPass00(PassManagerBuilder::EP_EnabledOnOptLevel0, registerPass);
```

```

class IntWrite: public FunctionPass {
public:
    static char ID; // Pass identification, replacement for typeid
    IntWrite() : FunctionPass(ID) {}

    Type *intTy, *ptrTy, *voidTy, *boolTy ; // To store the type instances
    FunctionCallee p_init ; // points to the function instance of _init_.
    FunctionCallee p_probe ; // points to the function instance of _probe_.

    IRBuilder<> * IRB;

    virtual bool doInitialization(Module &M) {
        /* doInitialization() is executed once per target module,
         * and executed before any invocation of runOnFunction().
         * This function is for initialization and the module level
         * instrumentation (e.g., add functions). */
        ...
        return true ; }

    virtual bool doFinalization(Module &M) {
        /* This function is executed once per target module after
         * all executions of runOnFunction() under the module. */

        return false ;}

    virtual bool runOnFunction(Function &F) {
        /* This function is invoked once for every function in the target
         * module by LLVM */

        /* Invoke runOnBasicBlock() for each basic block under F. */
        for (Function::iterator itr = F.begin() ; itr != F.end() ; itr++) {
            runOnBasicBlock(*itr) ;
        }

        return true;    }

    bool runOnBasicBlock (BasicBlock &B) {
        /* This function is invoked by runOnFunction() for each basic block
         * in the function. Note that this is not invoked by LLVM and different
         * from runOnBasicBlock() of BasicBlockPass.*/
        ...
        return true ; }

};} ...

```

More Information

- Writing an LLVM Pass
 - <http://llvm.org/docs/WritingAnLLVMPass.html>
- LLVM API Documentation
 - <http://llvm.org/doxygen/>
- How to Build and Run an LLVM Pass for Homework#4
 - <https://swtv.kaist.ac.kr/courses/cs453-fall14>